# Simple phase plane analysis and parameter estimation in R

*Rob J de Boer, Theoretical Biology, Utrecht University, 2018*

Phase plane analysis is powerful graphical method to analyze low-dimensional ODE models. At Utrecht University we have hitherto used dedicated C-code (GRIND) for numerical integration, phase plane analysis, and stability analysis of steady states. Thanks to the R-packages `deSolve` and `rootSolve` developed by Karline Soetaert and colleagues [1, 2, 4], it was relatively easy to copy most of GRIND's capabilities into R. People liking R may also like this simple interface to phase plane analysis. Thanks to the `FME` package, also developed by Karline Soetaert and colleagues [3], it was also possible to extend GRIND with non-linear parameter estimation using a least-squares approach. This resulted in an R-script `grind.R` defining five easy-to-use functions:

- `run()` integrates a model numerically and provides a time plot or a trajectory in the phase plane,
- `plane()` draws nullclines and can provide a vector field or phase portrait,
- `newton()` finds steady states (using the Newton-Raphson method) and can provide the Jacobian with its eigenvalues and eigenvectors.
- `continue()` performs parameter continuation of a steady state, providing a bifurcation diagram,
- `fit()` fits a model to data by estimating its parameters, and depicts the result in a timeplot.

The `run()` function calls `ode()` from the `deSolve` library, the `fit()` function calls `modFit()` from the `FME` library, and `newton()` and `continue()` call `steady()` from the `rootSolve` library. For the library functions one can get more help by typing `?...`, where ... is the name of the function (e.g., `?ode`). One can get help on the Grind functions by typing `args(...)`, where ... is one of the five Grind functions. The following sections are tutorials illustrating the usage of Grind. Instructions for installation are given in the final section.

# 1 Phase plane analysis

**Lotka Volterra model**. The ODEs of the model are defined in the simple notation defined for the `deSolve` package. The following is an example of the Lotka Volterra model, here defined by the function `model()`. This R-script is available as the file `lotka.R` on the website http://tbb.bio.uu.nl/rdb/grindR/:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
    dR <- r*R*(1 - R/K) - a*R*N
    dN <- c*a*R*N - delta*N
    return(list(c(dR, dN)))
  })
}
p <- c(r=1,K=1,a=1,c=1,delta=0.5) # p is a named vector of parameters
s <- c(R=1,N=0.01)                # s is the state
```

where the two lines below the function define the parameter values in the vector `p`, and the initial state of the variables in the vector `s`. Note that the function needs a current `state` and returns a list of derivatives (`dR,dN`). Importantly, *the derivatives should be specified in the same order as the variables in the state*. The names `model`, `s`, and `p` are the default designations for the model, state, and parameter values in all Grind functions. This example should be self explanatory as it just defines the Lotka Volterra model $dR/dt = rR(1 - R/K) - aRN$, and $dN/dt = caRN - \delta N$, with its parameter values and initial state as R-vectors, `p <- c(r=1,K=1,a=1,c=1,d=1,delta=0.5)`, and `s <- c(R=1,N=0.01)`, respectively. A shiny interactive primer on this model is available on the website: https://grind.shinyapps.io/lotka/.

The following tutorial is an example session illustrating the usage of the first four Grind functions, by simulating and analyzing this Lotka Volterra model (see Fig. 1 for its graphical output):
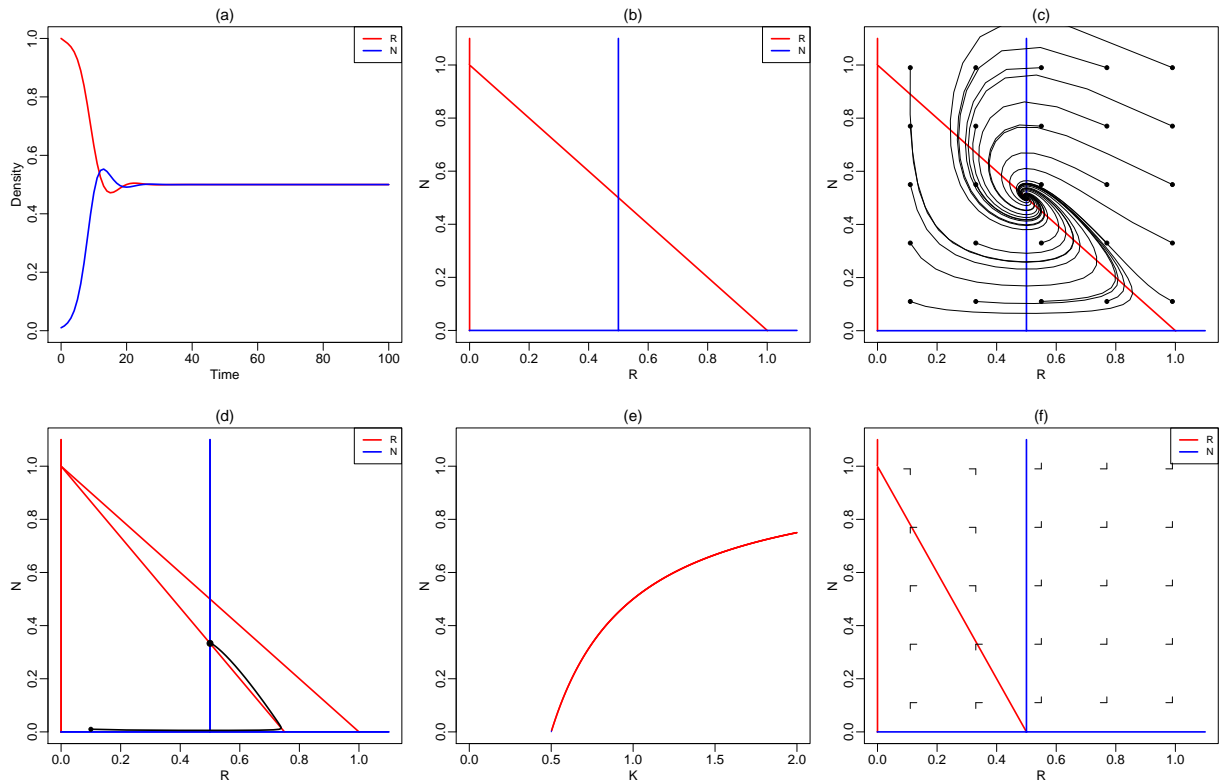
Figure 1: Numerical integration, phase plane analysis, and a bifurcation diagram of the Lotka Volterra model. The six panels collect the graphical output of the example session listed above.

```
run()                                   # run the model and make a time plot       (Fig 1a)
plane()                                 # make a phase plane with nullclines
plane(xmin=-0.001,ymin=-0.001)          # include the full axis in the phase plane  (Fig 1b)
plane(tstep=0.5,portrait=T)             # make a phase portrait                     (Fig 1c)
plane()                                 # make a clean phase plain again            (Fig 1d)
p["K"] <- 0.75                          # change the parameter K from 1 to 0.75
plane(add=T)                            # add the new nullclines
s["R"] <- 0.1                           # change the initial state to (R=0.1,N=0.01)
run(traject=T)                          # run the model and plot a trajectory
newton(c(R=0.5,N=0.5),plot=T)           # find a steady state around (R=0.5,N=0.5)a  (Fig 1d)
f <- newton(c(R=0.5,N=0.5))             # store this steady state in f
continue(f,x="K",xmax=2,y="N")          # continue this steady state while varying K (Fig 1e)
continue(f,x="K",xmax=2,y="N",step=0.001)  # get a better value with a smaller step size
p["K"] <- 0.5                           # set K to the value at which N goes extinct
plane(vector=T)                         # make a phase plane for this value of K     (Fig 1f)
```

Because all optional arguments to the Grind functions are local, subsequents calls to these functions do not adopt the previous definition of these options. For instance, if the axes are redefined by `plane(x=2,y=1)` and one wants to add a trajectory, one has to repeat the axis arguments, i.e., `run(x=2,y=1,traject=TRUE)`. This even hold for adding nullclines to an existing phase plane, e.g., `plane(xmax=2,add=TRUE)`. Defining simple shortcuts around the Grind functions allow you to implement your own preferences, e.g.,

```
pl <- function(...)  plane(x=2,y=1,eps=-0.001,...),
tr <- function(...)  run(x=2,y=1,traject=TRUE,...).
```

**Lac-operon model**. A slightly more sophisticated example shows how one can continue steady states to make a bifurcation diagram with a saddle-node bifurcation. We use a Lac-operon model defined in the reader of our Systems Biology course at Utrecht University. The graphical output of this example is displayed in Fig. 2, where bullets indicate stable steady states and circles depict unstable equilibria. We start close to the three steady state, call the Newton-Raphson algorithm, and store these states in the variables `low`, `mid`, and `hig`, respectively. Using parameter continuation we make a bifurcation

Figure 2: Nullclines and a bifurcation diagram of the Lac-operon model.

diagram in which we follow the middle steady state as a function of the external lactose concentration, $L$, by a call to `continue(mid,...)`. This R-script is available as the file `operon.R`:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
  R = 1/(1+A^n)                 # Repressor
  dA = M*L - delta*A - v*M*A  # Allolactose
  dM = c0 + c*(1-R) - d*M      # mRNA
  return(list(c(dA, dM)))
  })
}


p <- c(L=1,c=1,c0=0.05,d=1,delta=0.2,n=5,v=0.25)
s <- c(A=0,M=0)
plane(xmax=4)
low <- newton(s,plot=T)
mid <- newton(c(A=0.8,M=0.2),plot=T)
hig <- newton(c(A=2,M=1),plot=T)
continue(mid,x="L",y="A",xmax=2,ymax=4)
```

These two tutorials should be a sufficient introduction for standard phase plane analyses. The following sections illustrate the usage of events and noise (Section 2), vectors of equations (Section 3), and parameter estimation (Section 4), and can be read when needed. Section 5 is a reference manual that can be skimmed through and consulted when required. Finally, Section 6 provides simple installation instructions.

## 2 Combining numerical integration with events

The `deSolve` package allows one to execute discrete events while integrating the model numerically by using the `events` argument (see the `ode()` manual). This remains possible in Grind because `run()` passes additional options to `ode()` via the ellipsis (`...`) argument in R. We have added a somewhat simpler option (`after`) to handle events that are executed after each time step within `run()`. For instance `run(after="state<-ifelse(state<1e-9,0,state)")` will set small variables to zero after each time step. One can also stop when any of the variables becomes negative by `after <- "if(min(state)<0)break"`. This new option is illustrated by the following three examples each providing an R-command as a text, `after="text"` in a call of `run()` of the Lotka Volterra model introduced above.

For example, `after="parms[\"r\"]<-rnorm(1,1,0.1)"` sets the parameter $r$ to a random value, drawn from a normal distribution with a mean of one and a standard deviation of 0.1 (see the result in Fig. 3a). Note that `p` is called `parms` within `run()` (see the Manual), and the back-slashes in `\"r\"` before the quotes around the parameter name, because these quotes would otherwise mark the beginning or ending of a text. (Since $r$ is the first parameter, one can also just write

3

Figure 3: The three runs with `after="text"` executed after every timestep

`"parms[1]<-rnorm(1,1,0.1)"` to achieve the same effect). This random resetting of $r$ is done every timestep (as defined by the parameter `tstep=1` in `run()`).

The second example, `run(after="if(t==20)state[\"N\"]<-0")`, sets the predators $N = 0$ when time $t = 20$ (see Fig. 3b). Note again that `s` is called `state` in `run()` (see the Manual), and the backslashes in `\"N\"`. Again, the more simple `state[2]<-0` would achieve the same effect, because $N$ is the second variable. Finally, the third example adds Gaussian noise to both variables, e.g., `after="state<-state+rnorm(2,0,0.01)"` (see Fig. 3c). Note that `rnorm(2,0,0.01)` provides two random values, that are added to the two variables, respectively. The integration starts close to the steady state to prevent problems arising from random values setting a population to a negative value. These three examples are combined in the R-script `events.R`:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
    dR <- r*R*(1 - R/K) - a*R*N
    dN <- c*a*R*N - delta*N
    return(list(c(dR, dN)))
  })
}

p <- c(r=1,K=1,a=1,c=1,delta=0.5)
s <- c(R=1,N=0.01)

run(after="parms[\"r\"]<-rnorm(1,mean=1,sd=0.1)")
run(after="if(t==20)state[\"N\"]<-0")

# Use arrest to handle events at time points within time steps:
run(50,arrest=33.14,after="if(t==33.14)state[\"N\"]<-0",table=T)

f <- newton(c(R=0.5,N=0.5))
run(state=f,after="state<-state+rnorm(2,mean=0,sd=0.01)",ymax=1)

# Here is an example of similar event handling in deSolve:
fun<-function(t, y, parms){y["N"]<-0;return(y)}
run(events=list(func=fun,time=20))
```

**A final "event" to tweak the numerical data**. To modify the numerical solution data computed within `run()` one can pass on any R-command as a text with the `tweak` option. For instance, `run(tweak="nsol$T<-nsol[,2]+nsol[,3]")`, adds a fourth column to the solution by summing the first and second variable, and calls this column "T" (for total). Note that the numerical solution is called `nsol`, and that the first column contains the time. This manipulated `nsol` table is subsequently passed on to `timePlot`, or printed to screen (with the `table=TRUE` option in `run()`). This `tweak` option can be very helpful when fitting data containing columns representing (transformed) combinations of the variables of the model. Finally, note that this same tweak can be done by using apply to call sum: `run(tweak="nsol$T<-apply(nsol[,2:3],1,sum))"`, which shows how one can generalize
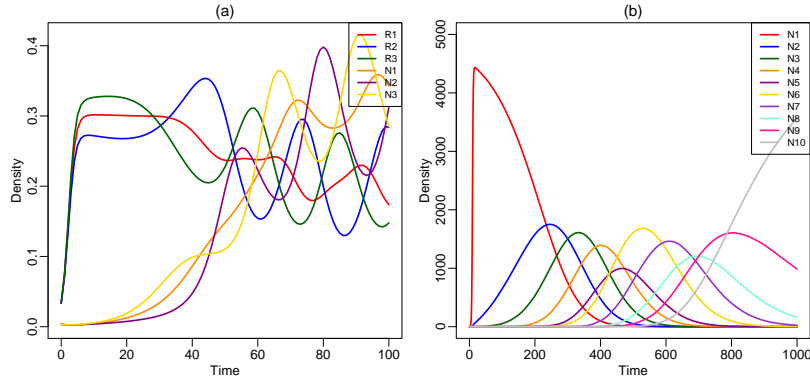
4

Figure 4: Left: the graphical output of the "vector of equations" example. Right: the output of the examples of mutating strains.

this to sum over a subset of columns.

# 3 Delay differential equations and maps

One can study maps by simply switching to Euler integration:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
    dN <- r*N*(1 - N) - N
    return(list(c(dN)))
  })
}
p <- c(r=3.75)
s <- c(N=0.01)
data <- run(1000,method="euler",table=TRUE)
plot(data$N[1:999],data$N[2:1000],pch=".")
```

Be careful with using the steady state functions (`newton()`, `continue()`, and `plane()`) because they are not aware of the fact that this model defines a map. Steady state values of this model will be correct because of the `-N` in the equation (i.e., `dN=0` has the same solution as $N_{t+1} = rN_t(1 - N_t)$), but the reported eigenvalues no longer define the stability of steady states.

One can study delay differential equations (DDEs) because the `deSolve` package implements a general solver (`dede()`) using the same syntax as the general ODE solver (`ode()`). For instance, the Lotka Volterra model with delayed growth of the predator,

$$\frac{\mathrm{d}R(t)}{\mathrm{d}t} = rR(t)(1 - R(t)/K) - aR(t)N(t) \ , \quad \text{and} \quad \frac{\mathrm{d}N(t)}{\mathrm{d}t} = caR(t - \Delta)N(t - \Delta) - \delta N(t) \ ,$$

would look like:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
    tlag <- t - Delta
    if (tlag < 0) lags <- c(0,0)   # no initial predation
    else lags <- lagvalue(tlag)    # returns lags of R and N
    dR <- r*R*(1 - R/K) - a*R*N
    dN <- a*lags[1]*lags[2] - d*N
    return(list(c(dR, dN)))
  })
}

p <- c(r=1,K=1,a=1,c=1,d=0.5,Delta=10)
s <- c(R=1,N=0.1)
run(delay=TRUE)
```

5

where the option `delay` tells Grind to use `dede()` solver. This would correspond to the situation where a prey at carrying capacity starts to be eaten by a predator at time zero, but the predator will only start to grow `Delta` time steps later. The `deSolve` function `lagvalue()` stores previous values of $R$ and $N$ in a vector (that is called `lags[]` here) that can be indexed to obtain the time lagged $R$ and $N$ values. Read the `deSolve` manual for further documentation and/or type `?dede` for help. Please note that the steady state functions (`newton()`, `continue()`, and `plane()`) should not be called with models containing calls to `lagvalue()`. Fitting DDE models to data should work fine (but solving DDEs is notoriously tricky).

# 4 Vectors of equations

Here is an example of a model with $n = 3$ prey populations, $R_i$, that are competing with each other via a logistic term (see the left panel in Fig. 4). Each prey has its own predator $N_i$,

$$\frac{\mathrm{d}R_i}{\mathrm{d}t} = b_i R_i (1 - \sum R_i) - d_1 R_i - a R_i N_i \quad \text{and} \quad \frac{\mathrm{d}N_i}{\mathrm{d}t} = a R_i N_i - d_2 N_i \ .$$

In the R-script `vector.R` we draw random prey birth rates, $b_i$, from a normal distribution:

```
model <- function(t, state, parms){
  with(as.list(c(state,parms)),{
    R <- state[1:n]
    N <- state[(n+1):(2*n)]
    S <- sum(R)
    dR <- b*R*(1-S) - d1*R - a*R*N
    dN <- a*R*N - d2*N
    return(list(c(dR,dN)))
  })
}
```

```
n <- 3                                  # number of species
b <- rnorm(n,mean=1,sd=0.1)             # b is a global parameter
p <- c(d1=0.1,d2=0.2,a=1)               # other parameters
R <- rep(0.1/n,n)                       # initial condition of R
names(R) <- paste("R",seq(1,n),sep="")  # Name them R1, R2, ... Rn
N <- rep(0.01/n,n)                      # initial condition of N
names(N) <- paste("N",seq(1,n),sep="")  # Name them N1, N2, ... Nn
s <- c(R,N)                             # combine R and N into s
run()
```

**Mutations after each time step**. Combining vectors and events one can model a series of evolving bacterial strains with increasing replication rates (see the right panel in Fig. 4). Consider the following model,

$$\frac{\mathrm{d}N_i}{\mathrm{d}t} = b_i N_i (1 - \sum N_i) - d N_i \ , \quad \text{for} \quad i = 1, 2, \ldots, n$$

and let strain $N_{i+1}$ evolve from strain $N_i$ at a mutation rate $\mu$. When the expected number of mutants, $\mu N_i$, is smaller than a single bacterium we calculate the probability that a single mutant appears, and add a single cell to $N_{i+1}$ (and subtract it from $N_i$). Otherwise the expected number of mutants is added to strain $N_{i+1}$ (and subtracted from $N_i$). This is realized by using `after` to call the function `mutate()`. In the R-script `evolve.R` birth rates $b_i$ of the strains increase linearly from $b_i = 1$ to $b_n = 2$:

```
model <- function(t, state, parms){
  with(as.list(c(state,parms)),{
    N <- state[1:n]
    S <- sum(N)
    dN <- b*N*(1-S/K) - d*N
    return(list(c(dN)))
  })
}
```
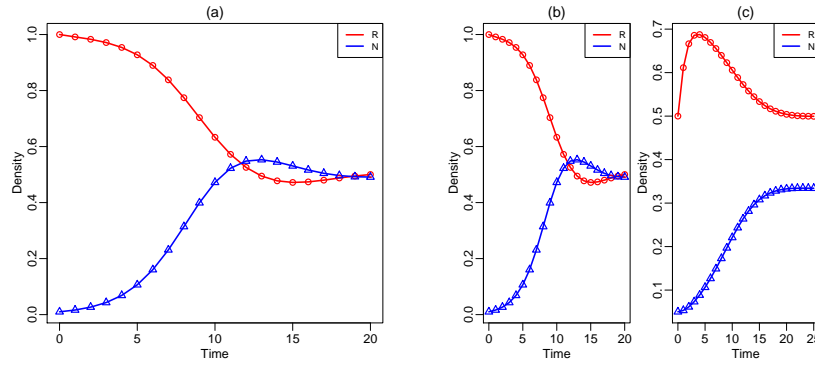
Figure 5: Fitting the Lotka Volterra model to data. Lines show the model for the best estimated parameters, and the symbols depict the data.

```
mutate <- function(t, state, parms){
  nmut <- rep(0,n+1)
  emut <- parms["mu"]*state                              # Expected number of mutants
  nmut[2:(n+1)] <- ifelse(emut>1,emut,ifelse(runif(n)<emut,1,0))
  state <- state + nmut[1:n] - nmut[2:(n+1)]
  return(state)
}

n <- 10                                                  # number of variants with
b <- seq(1,2,length=n)                                   # increasing birth rates
p <- c(d=0.1,K=5000,mu=0.001)                            # other parameters
s <- rep(0,n)                                            # set all variables to 0
s[1] <- 1                                                # set first to 1
names(s) <- paste("N",seq(1,n),sep="")                   # add names
run(1000,ymax=5000,after="state<-mutate(t,state,parms)")
```

# 5  Parameter estimation

One can fit the parameters of a model to data using the function `fit()`. When this function is called without any options it is assumed that there is `data.frame` called `data` with column names corresponding to the variables of the model, and `fit()` will then use the state, `s`, and all parameters, `p`, as an initial guess for fitting the data. We illustrate this using the same Lotka Volterra model as above, after creating a data set using `run(table=T)`. In the example we randomize the initial condition, `s`, and parameters, `p`, using a normal distribution (`rnorm()`) before we fit the model to the data (see Fig. 5A).

`fit()` internally calls the function `modFit` from the `FME` package (use `?modFit` to see all options), and delivers an object providing the estimated parameters, confidence ranges, and correlations between the parameters. Calling `fit()$par` just returns the estimated parameters (see below). Use the option `free` to explicitly define which parameters are "free" and should be estimated (`free` is a vector of names). The boolean option `initial` can be used to read the initial condition from the data (instead of estimating it).

One can fit simultaneously several data sets by provide a list of data sets to the first `datas` option (see the `fit(list(dataR,dataN))` example). When fitting several data sets, some of the parameters could be the same across all data sets, whereas others could differ, and have a unique value in each data set (see Fig. 5B & C). The shared parameters remain to be provided by the `free` option, and the unique parameters are provided by the `differ` option (see the `fit(list(data,data2),differ=c("R","N","K"))` example below). Here `differ` is just a vector of parameter names. However, in case one needs to supply an initial guess for the different parameters in each data set, one should define `differ` as a named list, containing the various guesses (see the `differ<-list()` example below). When data sets

differ in parameters that are known, these parameter values can be provided with the option `fixed`, which also has to be a list (see the lines below `fixed<-list()` in the example). Finally, we show how one can bootstrap the data by sampling (with recruitment) from every individual data set, and re-fit the samples using the best parameters as an initial guess. The following script is available as the file `lotka_fit.R`:

```
model <- function(t, state, parms) {
  with(as.list(c(state,parms)), {
    dR <- r*R*(1 - R/K) - a*R*N
    dN <- c*a*R*N - delta*N
    return(list(c(dR, dN)))
  })
}


p <- c(r=1,K=1,a=1,c=1,delta=0.5)
s <- c(R=1,N=0.01)
data <- run(20,table=T)          # Make a data set

s <- s*abs(rnorm(2,1,0.1));s    # Random guess for initial condition
p <- p*abs(rnorm(5,1,0.1));p    # Random guess for parameters
f <- fit()                      # Fit data with all 7 parameters free
summary(f)                      # Check confidence ranges, etcetera
p <- f$par[3:7];p               # Store estimates in p

p <- p*abs(rnorm(5,1,0.1));p    # Another random guess for the parameters
w <- c(names(s),names(p))       # w provides the names of free parameters
f <- fit(data,free=w)           # Fit the data again
f <- fit(data,initial=T,free=names(p))   # Take initial condition from data

dataR <- data; dataR$N <- NULL # Make two data sets one with R,
dataN <- data; dataN$R <- NULL # and the other with N,
f <- fit(list(dataR,dataN))    # which gives the same result

p <- c(r=1,K=1,a=1,c=1,delta=0.5)  # Start again with same parameters
p["K"] <- 0.75                     # Change K,
s <- c(R=0.5,N=0.05)               # and the initial condition,
data2 <- run(25,table=T)           # and make a new data set
s <- c(R=0.75,N=0.02)              # An "average" guess for the 2 initial conditions
par(mfrow=c(1,2))                  # Show two panels next to each other
f <- fit(list(data,data2),differ=c("R","N","K"),main=c("A","B"))
f$par                              # Show parameters only

# Provide individual initial guesses as a list:
differ <- list(R=c(0.9,0.55),N=c(0.02,0.04),K=c(1.1,0.7))
f <- fit(list(data,data2),differ=differ,main=c("A","B"))

# Provide fixed parameters as a list:
fixed <- list(R=c(1,0.5),N=c(0.01,0.05))
differ <- "K"                             # one unknown parameter (K)
free <- names(p)[-2];free                 # and four shared unknown parameters
f <- fit(list(data,data2),free=free,differ=differ,fixed=fixed,main=c("A","B"))

# The latter is identical to taking the initial condition from the data:
f <- fit(list(data,data2),free=free,differ=differ,initial=T,main=c("A","B"))

# Finally perform a 100 bootstrap simulations:
fit(list(data,data2),free=free,differ=differ,fixed=fixed,main=c("A","B"),bootstrap=100)$par
par(mfrow=c(1,1))
```

# 6    Manual

A model can be solved numerically from the initial state by calling `run()`, and the output will be a timeplot, trajectory or table. Next to the graphics output, `run()` returns the final state attained by the simulation (or all data when `table=TRUE`). The former can be helpful if one wants to continue from the previous state (e.g., `f<-run(); f<-run(state=f)`). The full definition of `run()` is:

```
run <- function(tmax=100, tstep=1, state=s, parms=p, odes=model, ymin=0, ymax=NULL,
log="", x=1, y=2, xlab="Time", ylab="Density", tmin=0, draw=lines, times=NULL,
show=NULL, arrest=NULL, after=NULL, tweak=NULL, timeplot=TRUE, traject=FALSE,
table=FALSE, add=FALSE, legend=TRUE, solution=FALSE, delay=FALSE, lwd=2,
col="black", pch=20, ...)
```

`run()` calls the `ode()` function from the `deSolve` package. Additional arguments (...) are passed on to `ode()` and `plot()`.

The phase plane function `plane()` sets up a space with the first variable on the horizontal axis, and the second on the vertical axis. The full definition of `plane()` is:

```
plane <- function(xmin=0, xmax=1.1, ymin=0, ymax=1.1, log="", npixels=500, state=s,
parms=p, odes=model, x=1, y=2, time=0, grid=5, eps=0, show=NULL, portrait=FALSE,
vector=FALSE, add=FALSE, legend=TRUE, zero=TRUE, lwd=2, col="black", pch=20, ...)
```

Additional arguments (...) are passed on to `run()` (for the phase portrait) and to `plot()`. Note that `plane()` calls the "vectorized" R-function `outer()`, which implies that if one calls functions in the ODEs they should also be vectorized, e.g., one should use `pmax()` instead of `max()`. Finally, note that there is an extension, `cube.R`, for 3-dimensional nullclines.

The function `newton()` finds a steady state from a nearby initial state, and can report the Jacobi matrix with its eigenvalues and eigenvectors. The full definition of `newton()` is:

```
newton <- function(state=s, parms=p, odes=model, time=0, x=1, y=2, positive=FALSE,
jacobian=FALSE, vector=FALSE, plot=FALSE, ...)
```

`newton()` calls the function `steady()` from the `rootSolve` package (which calls `stode()`). Additional arguments (...) are passed on to both of them. `newton()` needs an initial state close to an equilibrium point.

The function `continue()` continues a steady state by changing a "bifurcation" parameter defined by the horizontal axis of the bifurcation diagram. The full definition of `continue()` is:

```
continue <- function(state=s, parms=p, odes=model, x=1, step=0.01, xmin=0, xmax=1,
y=2, ymin=0, ymax=1.1, log="", time=0, positive=FALSE, add=FALSE, ...)
```

`continue()` calls the function `steady()`from the `rootSolve` package (additional arguments (...) are passed on), and needs an initial state close to an equilibrium point. Note that there is much more proper software for bifurcation analysis like XPPAUT or MatCont, which reports the type of bifurcations encountered, and automatically continues all branches of branch points.

The function `fit()` fits a model to data by non-linear parameter estimation. The output is an object (class of `modFit`) containing the estimated parameters, the summed squared residuals, confidence ranges, and correlations (see the `modFit()` manual). The data and the model behavior for its best fit parameters are depicted in a timeplot. Its full definition is:

```
fit <- function(datas=data, state=s, parms=p, odes=model, free=NULL, differ=NULL,
fixed=NULL, tmin=0, tmax=NULL, ymin=0, ymax=NULL, log="", xlab="Time",
ylab="Density", bootstrap=0, show=NULL, fun=NULL, costfun=cost, initial=FALSE,
add=FALSE, timeplot=TRUE, legend=TRUE, main=NULL, sub=NULL, pchMap=NULL, ...)
```

`fit()` calls the function `modFit()` from the `FME` package (which calls `modCost()`). Additional arguments (...) are passed on to both of them, and to `run()` and `ode()`.

Finally the internal function `timePlot()` can be used to plot data and is defined as:

```
timePlot <- function(data, tmin=0, tmax=NULL, ymin=0, ymax=NULL, log="",
xlab="Time", ylab="Density", show=NULL, legend=TRUE, draw=lines, lwd=2, add=FALSE,
main=NULL, sub=NULL, colMap=NULL, pchMap=NULL, ...)
```

These functions have many arguments, and fortunately most of them have good default values, and can hence typically be omitted. The arguments can be used to define various options and adjustments:

- `state=s, parms=p, odes=model` define the names of the state vector, parameter vector, and the model.
- `tmax=100, tstep=1` set the integration time and the reporting interval. `tmin` allows one to start a specific time point (which can be convenient when a run is continued). One can also provide a vector of time points where the `run()` should provide output with the `times` option (see the `ode()` manual).
- `x=1, y=2` define the variables on the axes of phase planes and bifurcation diagrams. One can also use the names of the variables to define the axes, e.g., `x="R",y="N"`.
- `xmin=0, xmax=1.1, ymin=0, ymax=1.1, log=""` define the scaling of the horizontal and vertical axes of phase planes and bifurcation diagrams (`log="y"` makes the vertical axis logarithmic).
- `step=0.01` defines the maximum change of the bifurcation parameter in a bifurcation diagram. When the axis is linear the parameter is increased, or decreased, in steps not exceeding `step` × `xmax`. When the axis is logarithmic the parameters is maximally multiplied by 1+`step`. `continue()` will decrease the step size to maximally `step`/100 when it looses the steady state.
- `xlab="Time", ylab="Density"` allow one to redefine the labels of the axes of a time plot.
- `show=NULL` defines the variables appearing in a time plot, fit, or phase plane. By default all are shown. By explicitly providing a list of variables, one can omit some of the variables. `show` is typically a list of names (`show=c("P","Q")`).
- `after=NULL` defines a command to be executed after each time step, e.g., `after="state <- ifelse( state<1e-9, 0, state)"` sets small variables to zero.
- `arrest=NULL` defines a vector of values, or parameter names, defining time points where the integrator should stop, and report the current state (i.e., these time points are added to the `times` vector of `ode()`). This can be helpful when fitting a piece-wise model for which the discontinuous time points have to be estimated (e.g., `arrest=c("T1","T2")`).
- `tweak=NULL` allows one to modify the data delivered by `run()`. For instance one can add columns that can be fitted to data: (`tweak<-"nsol<-cbind(nsol,nsol[,2]+nsol[,3]);names(nsol)[4]<-\"T\""`), or transform the simulation data before they are fitted to data that is already transformed.
- `timeplot=TRUE, traject=FALSE, table=FALSE` determine the output of `run()` in the form of a time plot (default), trajectory in a phase plane, and/or a table with all data.
- `draw=lines` draws the timeplot as continuous lines. The alternative is `draw=points`.
- `lwd=2` sets the line width of the graphs, `colMap=NULL` and `pchMap=NULL` can be used to re-map the colors or symbols, e.g., `pchMap=c(3,2,1)` reverts the order of the first three R-symbols (see `pch` in `points()`). Note that `grind.R` defines its own color table (of dark colors that print well). If you don't like this, uncomment the second `colors <-` line in the `grind.R` script.
- `main=NULL, sub=NULL` allow one to put a title at the top and/or a subtitle at the bottom of the graph (these are passed on to the R-function `plot()`). Note that these titles are set in a plain font (to set that back to bold face, change the two `font` lines in the `grind.R` file).
- `add=FALSE, legend=TRUE` define whether or not the new plot should be added to the current one, and a legend should be placed.
- `solution=FALSE` tells `grind.R` whether or not the model provides time derivatives (default), or a full solution, which is particularly useful when fitting data to functions. This can only be used in `run()` and `fit` and should obviously not be used in combination with phase plane analysis, nor with searching steady states (`newton()`, `continue()`). Models returning a solution obey the same format as the ODE models required by `deSolve`, except for the fact that they should return a value, or a vector of values (and not a list).
- `delay=FALSE` tells `grind.R` whether or not to call the DDE solver from the `deSolve` package. The time delay(s) in the model are to be defined by calling the `lagvalue()` function from the `deSolve` package. This can only be used in `run()` and `fit` and should also not be used in combination with

phase plane analysis, nor with searching steady states (`newton()`, `continue()`).

- `npixels=500` defines the resolution of the phase space in `plane()`
- `time=0` defines the time point for which nullclines are computed and steady states are computed (for non-autonomous ODEs).
- `grid=5` defines the number of grid points for which the vector field or phase portrait is drawn.
- `eps=0` is a shortcut in `plane()` to include or exclude the axis in the nullclines: `eps` is added to both `xmin` and `ymin`.
- `portrait=FALSE`, `vector=FALSE` define whether or not `plane()` should include a phase portrait or vector field.
- `zero=TRUE` draws the phase plane for all variables other than `x` and `y` set to zero (also important when drawing nullclines of variables not appearing on the axes).
- `positive=FALSE`, setting `positive=TRUE` will restricts the search of `newton()` and `continue()` to positive steady states only.
- `jacobian=FALSE,vectors=FALSE,plot=FALSE` define whether or not `newton()` should print the Jacobian and eigenvectors, and indicate the steady state by a symbol in the phase plane.
- `datas=data` in `fit()` defines the name of the data frame containing the data, or defines a list of data frames.
- `free=NULL` defines the names of the parameters to be fitted. By default `free` equals `c(names(state), names(parms))`.
- `differ=NULL` defines the names of the parameters that differ between the data sets and need to be fitted separately. `differ` can also be a named list containing the individual guesses for each data set. (One can use `makelist(differ,state,parms,nsets)` to set up such a list).
- `fixed=NULL` defines a named list of the parameters that differ between the data sets and have known fixed values. (One can use `makelist(fixed,state,parms,nsets)` to set up such a list).
- `initial=FALSE` allows one to read the initial condition from the data (and not estimate it).
- `fun=NULL` defines a function to transform the data and (numerical) solution before fitting.
- `costfun=cost` allows one to redefine the cost-function measuring the distance between the model and the data. This can be useful when different data sets need different models. The default cost-function loops over the various data sets, and calls `run()` for each of them. That call can easily be adapted for each data set (the index of the loop is called `iset`).
- `bootstrap=0` defines the number of samples to be taken randomly from the data (with replacement). This prints a summary and adds an element `bootstrap` to the `modFit` list, containing a matrix with all parameter estimates. Use `pairs(f$bootstrap)` to see the correlations between the estimates.
- `...` can be used to define parameters that are passed on to other functions

# 7   Installation and startup

The very first time one may need to install the Soetaert libraries into the R-environment, e.g., `install.packages(c("deSolve", "rootSolve", "FME")` in R or Jupyter, or `Install Packages` in the `Tools` menu of RStudio. After downloading `grind.R` one can include the grind.R environment by typing `source("grind.R")`. In RStudio one can also open `grind.R` in one tab (and "source" it) and open the model in another tab. The R-scripts can be found on the webpage: `http://tbb.bio.uu.nl/rdb/grindR/`.

For example, download the R-codes `grind.R` and `lotka.R` in a local directory, and open them in RStudio (if RStudio is the default R-environment one can just double click the links). It may be convenient to set the working directory to the folder where the R-codes were stored (`Set working directory` in the `Session` menu of RStudio). First "source" the `grind.R` file (button in right hand top corner) to define the Grind functions. Then click the other tab, and select the function `model()` by highlighting everything up to the closing curly bracket, and execute this by clicking the `Run` button (or typing `Control Enter`). Subsequently proceed through that file by running it line-by-line (using `Control Enter`), to see what is happening, and become familiar with the behavior of the Grind

functions.

Once you have a picture that you like, you may copy the lines creating that figure into the `lotka.R` window for later usage. (Again, use "Run" or "Control Enter" to execute lines from the `lotka.R` panel into the console). Finally, one can reduce the amount of white space in the margins, or plot panels next to each other, by changing R's default graphics settings (`par()`); see the examples in `bbox.R`.

August 26, 2018, Rob J. de Boer

# References

[1] **Soetaert, K.**, 2009. rootSolve: Nonlinear root finding, equilibrium and steady-state analysis of ordinary differential equations. R package 1.6.

[2] **Soetaert, K. and Herman, P. M.**, 2009. A Practical Guide to Ecological Modelling. Using R as a Simulation Platform. Springer. ISBN 978-1-4020-8623-6.

[3] **Soetaert, K. and Petzoldt, T.**, 2010. Inverse modelling, sensitivity and Monte Carlo analysis in R using package FME. Journal of Statistical Software **33**:1–28.

[4] **Soetaert, K., Petzoldt, T., and Setzer, R. W.**, 2010. Solving differential equations in R: Package deSolve. Journal of Statistical Software **33**:1–25.

# A    Appendix

## A.1    What's in the box?

The files on the `http://tbb.bio.uu.nl/rdb/grindR/` website provide:

| | |
|---|---|
| `app.R` | . . . . . . . . example for making an app (see the website https://grind.shinyapps.io/lotka/) |
| `bbox.R` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . examples for setting the margins (using `par()`) |
| `chaos.R` | . . . an example of 3-dimensional nullclines (using `cube.R`), a Takens reconstruction of a chaotic attractor, and a bifurcation diagram formed by continuing the period doubling cascade of a limit cycle |
| `cube.R` | . . . . . . . . . an extension of Grind defining `cube()` and `run3d()` for plotting 3-dimensional nullclines and trajectories |
| `delay.R` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . an example of a DDE (delay differential equation) |
| `events.R` | . . . . . . . . . . . . . . . . . . . . . . . . . examples of event handling using the option `after` in `run()` |
| `evolve.R` | . . . . . . . . . . . . . . . an example using arrays of ODEs and implementing random mutations |
| `grind.R` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . the Grind wrapper around the Soetaert libraries |
| `grind2grindR.html` | . . . . . . . . . . . . . . . . a webpage translating previous GRIND commands into current Grind |
| `logisticMap.R` | . . . . . . the Logistic map, a Takens reconstruction of its chaotic attractor, and its famous bifurcation diagram |
| `lotka.R` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . the Lotka-Volterra model defined in the `deSolve` format |
| `lotka.ipynb` | . . . . . . . . . . . . . . . phase plane analysis of the Lotka-Volterra model in a jupyter notebook |
| `monod.R` | . . . . . . . . . . . . . . . a Monod saturated predator-prey model defined in the `deSolve` format |
| `operon.R` | . . . . . . . . . a model of the Lac-operon illustrating the usage of `continue()` revealing two saddle-node bifurcations |
| `README` | . . . . . . . . . . . . . . . . . . . . . a text file with some instructions and a history of Grind versions |
| `separatrix` | an example for drawing a stable manifold by starting on the corresponding eigenvector, and integrating backwards |
| `solution.R` | . . . . . . . . . . . . . . . . . an example of a `model()` returning a solution rather than derivatives |
| `tutorial.pdf` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . the Grind manual |
| `vector.R` | . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . an example with arrays of ODEs |

## A.2    Grind extensions

There is currently only one extension (`cube.R`) which uses the `plot3D` R-library to define a function for plotting 3-dimensional nullclines (`cube()`) and a function for plotting 3-dimensional trajectories (`run3d()`). The syntax of these functions is similar to that of `plane()` and `run()`, respectively. See the `cube.R` script for some help on the 3D projection.